

```
In [ ]: ## Multilinear Regresson
```

```
In [37]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

```
In [2]: from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler
from feature_engine.outliers import Winsorizer
```

```
In [3]: ## imoprt data
```

```
In [4]: data = pd.read_csv(r"C:\Users\shubham lokare\Downloads\Multiple_Linear Regression\Multiple_Linear Regression\Ca
```

```
In [5]: data
```

```
Out[5]:
```

	MPG	Enginetype	HP	VOL	SP	WT
0	53.700681	petrol	49	89	104.185353	28.762059
1	50.013401	hybrid	55	92	105.461264	30.466833
2	50.013401	diesel	55	92	105.461264	30.193597
3	45.696322	lpg	70	92	113.461264	30.632114
4	50.504232	petrol	53	92	104.461264	29.889149
...
76	36.900000	hybrid	322	50	169.598513	16.132947
77	19.197888	lpg	238	115	150.576579	37.923113
78	34.000000	hybrid	263	50	151.598513	15.769625
79	19.833733	diesel	295	119	167.944460	39.423099
80	12.101263	hybrid	236	107	139.840817	34.948615

81 rows × 6 columns

```
In [6]: ### check how many engine type
data['Enginetype'].unique()
```

```
Out[6]: array(['petrol', 'hybrid', 'diesel', 'lpg', 'cng'], dtype=object)
```

```
In [7]: data['Enginetype'].value_counts()
```

```
Out[7]: diesel    26
petrol    16
hybrid    16
lpg       12
cng       11
Name: Enginetype, dtype: int64
```

```
In [8]: ### check missing values
data.isna().sum()
```

```
Out[8]: MPG      0
Enginetype  0
HP          0
VOL         0
SP          0
WT          0
dtype: int64
```

```
In [9]: ### split the data into input and output variable
# Seperating input and output variables
X = pd.DataFrame(data.iloc[:, 1:6])
y = pd.DataFrame(data.iloc[:, 0])
```

```
In [10]: X    ### input variable
```

```
Out[10]:
```

	Enginetype	HP	VOL	SP	WT
0	petrol	49	89	104.185353	28.762059
1	hybrid	55	92	105.461264	30.466833
2	diesel	55	92	105.461264	30.193597
3	lpg	70	92	113.461264	30.632114
4	petrol	53	92	104.461264	29.889149
...
76	hybrid	322	50	169.598513	16.132947
77	lpg	238	115	150.576579	37.923113
78	hybrid	263	50	151.598513	15.769625
79	diesel	295	119	167.944460	39.423099
80	hybrid	236	107	139.840817	34.948615

81 rows × 5 columns

```
In [11]: y      ### output variable
```

```
Out[11]:
```

	MPG
0	53.700681
1	50.013401
2	50.013401
3	45.696322
4	50.504232
...	...
76	36.900000
77	19.197888
78	34.000000
79	19.833733
80	12.101263

81 rows × 1 columns

```
In [12]: ##### seprate the numerical and categorical columns

numerical_feature = X.select_dtypes(exclude=['object']).columns
numerical_feature
```

```
Out[12]: Index(['HP', 'VOL', 'SP', 'WT'], dtype='object')
```

```
In [13]: ### categorical data
categorical_feature = X.select_dtypes(include=['object']).columns
categorical_feature
```

```
Out[13]: Index(['Enginetype'], dtype='object')
```

```
In [14]: ##### make pipeline for missing values
numerical = Pipeline([('impute' , SimpleImputer(strategy = 'mean'))])
```

```
In [15]: ### transform into columns
process = ColumnTransformer([('impute' , numerical , numerical_feature)])
```

```
In [16]: data1 = process.fit(X)
```

```
In [17]: new_data = pd.DataFrame(data1.transform(X) , columns =numerical_feature)
```

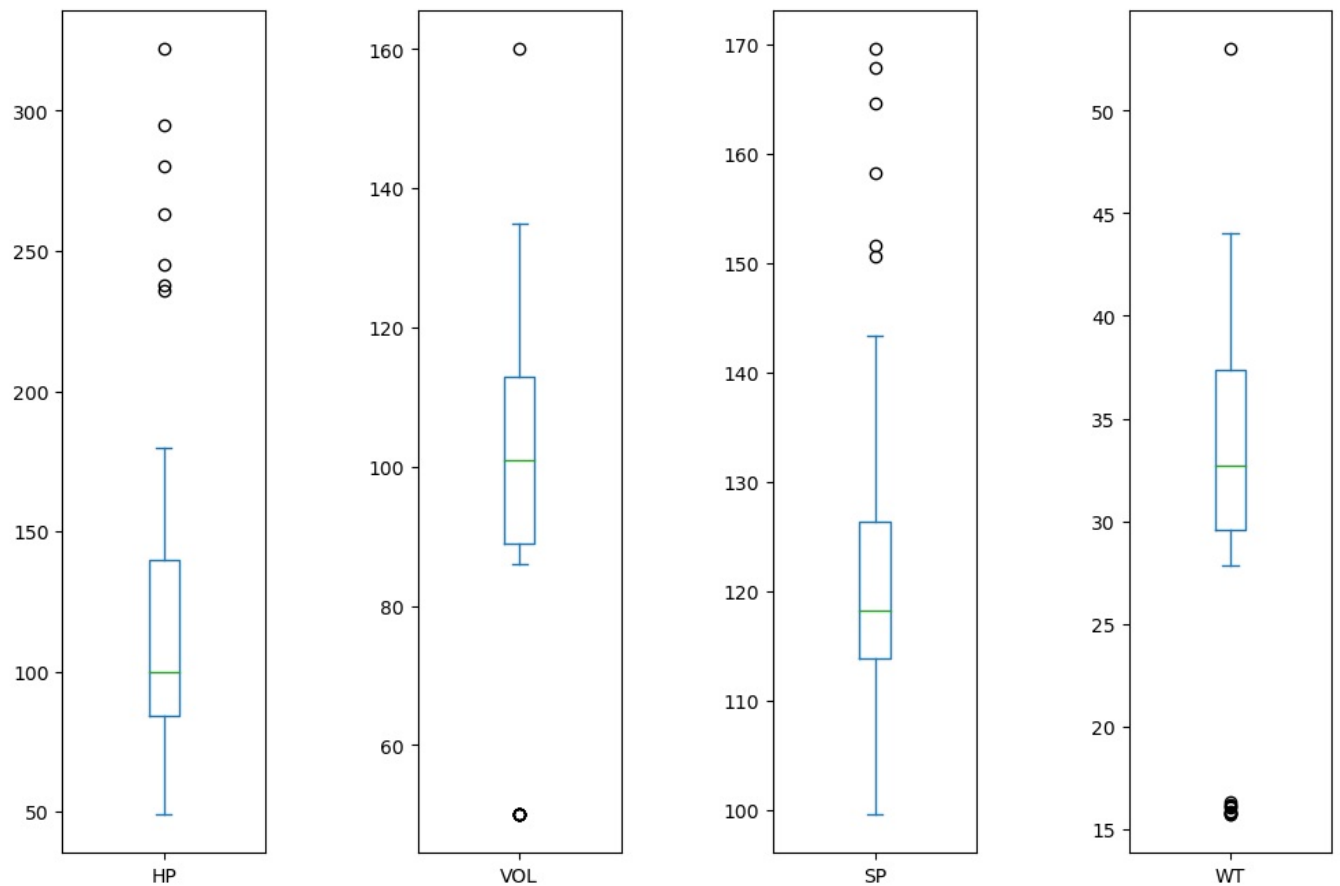
```
In [18]: new_data
```

```
Out[18]:
```

	HP	VOL	SP	WT
0	49.0	89.0	104.185353	28.762059
1	55.0	92.0	105.461264	30.466833
2	55.0	92.0	105.461264	30.193597
3	70.0	92.0	113.461264	30.632114
4	53.0	92.0	104.461264	29.889149
...
76	322.0	50.0	169.598513	16.132947
77	238.0	115.0	150.576579	37.923113
78	263.0	50.0	151.598513	15.769625
79	295.0	119.0	167.944460	39.423099
80	236.0	107.0	139.840817	34.948615

81 rows × 4 columns

```
In [19]: ### check outliers
X.plot(kind='box', subplots = True, figsize=(12,8))
plt.subplots_adjust(wspace= 0.75)
plt.show()
```




```
In [20]: # Winsorization for outlier treatment
winsor = Winsorizer(capping_method = 'iqr', # choose IQR rule boundaries or gaussian for mean and std
                    tail = 'both', # cap left, right or both tails
                    fold = 1.5,
                    variables = list(new_data.columns))
```

```
In [21]: clean = winsor.fit(new_data)
```

```
In [22]: new = pd.DataFrame(clean.transform(new_data), columns = numerical_feature)
```

```
In [23]: ### check outliers
new.plot(kind='box', subplots = True, figsize=(12,7))
```

```
Out[23]: HP      Axes(0.125,0.11;0.168478x0.77)
VOL      Axes(0.327174,0.11;0.168478x0.77)
SP       Axes(0.529348,0.11;0.168478x0.77)
WT       Axes(0.731522,0.11;0.168478x0.77)
dtype: object
```



```
In [24]: # Scaling
## Scaling with MinMaxScaler
scale = Pipeline([('scale', MinMaxScaler())])
```

```
In [25]: process3 = ColumnTransformer([('scale', scale, numerical_feature)])
# Skips the transformations for remaining columns
```

```
In [26]: data2 = process3.fit(new)
```

```
In [27]: clean_data = pd.DataFrame(data2.transform(new) , columns = numerical_feature)
```

```
In [28]: clean_data
```

```
Out[28]:
```

	HP	VOL	SP	WT
0	0.000000	0.375000	0.101099	0.348409
1	0.034286	0.406250	0.129017	0.403044
2	0.034286	0.406250	0.129017	0.394287
3	0.120000	0.406250	0.304063	0.408341
4	0.022857	0.406250	0.107136	0.384530
...
76	1.000000	0.000000	1.000000	0.000000
77	1.000000	0.645833	1.000000	0.642004
78	1.000000	0.000000	1.000000	0.000000
79	1.000000	0.687500	1.000000	0.690076
80	1.000000	0.562500	0.881269	0.546677

81 rows × 4 columns

```
In [29]: clean_data.describe()
```

```
Out[29]:
```

	HP	VOL	SP	WT
count	81.000000	81.000000	81.000000	81.000000
mean	0.369312	0.478781	0.456253	0.470733
std	0.266050	0.220992	0.244018	0.221019
min	0.000000	0.000000	0.000000	0.000000
25%	0.200000	0.375000	0.312113	0.375000
50%	0.291429	0.500000	0.407941	0.475719
75%	0.520000	0.625000	0.587268	0.625000
max	1.000000	1.000000	1.000000	1.000000

```
In [30]: ## Encoding
# Categorical features
encoding_pipeline = Pipeline([('onehot', OneHotEncoder())])
```

```
In [31]: preprocess_pipeline = ColumnTransformer([('categorical', encoding_pipeline, categorical_feature)])

clean = preprocess_pipeline.fit(X) # Works with categorical features only
```

```
In [34]: encode_data = pd.DataFrame(clean.transform(X).todense())

# To get feature names for Categorical columns after Onehotencoding
encode_data.columns = clean.get_feature_names_out(input_features = X.columns)
encode_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 81 entries, 0 to 80
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   categorical_EngineType_cng             81 non-null     float64
1   categorical_EngineType_diesel          81 non-null     float64
2   categorical_EngineType_hybrid          81 non-null     float64
3   categorical_EngineType_lpg             81 non-null     float64
4   categorical_EngineType_petrol          81 non-null     float64
dtypes: float64(5)
memory usage: 3.3 KB
```

```
In [35]: cleandata = pd.concat([clean_data, encode_data], axis = 1)
# concatenated data will have new sequential index
cleandata.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 81 entries, 0 to 80
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   HP                                     81 non-null     float64
1   VOL                                    81 non-null     float64
2   SP                                     81 non-null     float64
3   WT                                     81 non-null     float64
4   categorical_EngineType_cng             81 non-null     float64
5   categorical_EngineType_diesel          81 non-null     float64
6   categorical_EngineType_hybrid          81 non-null     float64
7   categorical_EngineType_lpg             81 non-null     float64
8   categorical_EngineType_petrol          81 non-null     float64
dtypes: float64(9)
memory usage: 5.8 KB
```

```
In [38]: ### plot pairplot
sns.pairplot(cleandata)
```

C:\Users\shubham lokare\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

```
self.figure.tight_layout(*args, **kwargs)
```

```
Out[38]: <seaborn.axisgrid.PairGrid at 0x1fb6210e490>
```



```
In [40]: ### # Library to call OLS model
# import statsmodels.api as sm

# Build a vanilla model on full dataset
from sklearn.model_selection import train_test_split
# import statsmodels.formula.api as smf
import statsmodels.api as sm

from sklearn.linear_model import LinearRegression
```

```
In [45]: X_train ,X_test , Y_train ,Y_test =train_test_split(cleandata ,y , test_size = 0.2 , random_state = 0)
```

```
In [48]: ## Build the best model Model building with OLS
model = sm.OLS(Y_train, X_train).fit()
model.summary()
```

Out[48]:

OLS Regression Results							
Dep. Variable:		MPG	R-squared:		0.835		
Model:		OLS	Adj. R-squared:		0.811		
Method:		Least Squares	F-statistic:		34.78		
Date:		Sun, 09 Jun 2024	Prob (F-statistic):		7.76e-19		
Time:		12:36:00	Log-Likelihood:		-174.70		
No. Observations:		64	AIC:		367.4		
Df Residuals:		55	BIC:		386.8		
Df Model:		8					
Covariance Type:		nonrobust					
		coef	std err	t	P> t	[0.025	0.975]
HP		-37.5602	8.292	-4.530	0.000	-54.177	-20.944
VOL		-25.3553	56.330	-0.450	0.654	-138.244	87.533
SP		14.6843	9.005	1.631	0.109	-3.362	32.731
WT		10.6664	56.722	0.188	0.852	-103.007	124.340
categorical__Enginetype_cng		47.9789	2.096	22.893	0.000	43.779	52.179
categorical__Enginetype_diesel		47.8560	2.082	22.986	0.000	43.684	52.028
categorical__Enginetype_hybrid		50.2603	1.882	26.701	0.000	46.488	54.033
categorical__Enginetype_lpg		48.6177	2.185	22.249	0.000	44.238	52.997
categorical__Enginetype_petrol		49.9100	1.893	26.369	0.000	46.117	53.703
Omnibus:		1.100	Durbin-Watson:		2.001		
Prob(Omnibus):		0.577	Jarque-Bera (JB):		0.878		
Skew:		0.286	Prob(JB):		0.645		
Kurtosis:		2.952	Cond. No.		166.		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [49]: pred = model.predict(X_test)

In [50]: pred

Out[50]: 22    37.018285
27    37.987385
61    31.025396
13    42.911360
71    21.751389
74    26.273910
30    37.306081
55    28.224419
53    24.718994
26    35.429283
50    32.267032
42    38.303991
48    34.342240
33    34.375215
73    24.121018
2     42.367794
57    32.226089
dtype: float64

In [53]: ## check r2score
from sklearn.metrics import r2_score

In [54]: score = r2_score(Y_test ,pred)

In [55]: score

Out[55]: 0.8027513141600667

In [56]: ## mean
np.mean(pred)

Out[56]: 32.97940477451849

In [57]: # Train residual values
error = Y_test.MPG -pred
```

```
In [58]: error
```

```
Out[58]: 22    1.292321
27    0.423618
61   -6.416264
13    1.741474
71    1.452180
74   -7.187569
30    2.125154
55   -0.368167
53   -0.231627
26    2.981720
50   -2.637096
42   -4.233323
48   -3.328109
33    1.910241
73   -5.034678
2     7.645607
57   -2.596153
dtype: float64
```

```
In [59]: ### ## Scores with Cross Validation (cv)
```

```
lm = lm = LinearRegression()
```

```
In [75]: from sklearn.model_selection import cross_val_score
```

```
cross = cross_val_score(lm ,X_train ,Y_train , scoring = 'neg_mean_squared_error' ,cv=8)
```

```
In [76]: cross
```

```
Out[76]: array([-28.49153209, -25.55480712, -13.76367249, -15.54257037,
        -18.35795527, -7.80334964, -40.0516394 , -5.1595768 ])
```

```
In [79]: ####check mean
np.mean(cross)
```


```
Out[79]: -19.340637897928936
```

```
In [83]: ### to check th model is good or not if variance is low then model is good
sns.displot(pred-Y_test.MPG ,kind='kde')
```

C:\Users\shubham lokare\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

self.figure.tight_layout(*args, **kwargs)

```
Out[83]: <seaborn.axisgrid.FacetGrid at 0x1fb06fa88d0>
```



```
In [ ]:
```